

Business Understanding

The objective was to use machine learning to predict customer churn for a telecom company. Understanding the churn problem helps SyriaTel maintain profitability and customer loyalty. A markdown section outlines this motivation and frames churn prediction as a classification task. I also explained the target variable and the potential business value of the solution. This gave context for the technical work that followed.

Preprocessing

I explored the dataset using descriptive statistics and visualizations to understand trends. I checked for missing values, class imbalance, and feature distributions. Categorical variables were encoded, and numerical features were scaled where needed. I split the data into training and test sets to avoid data leakage. This ensured a clean and fair modeling process.

In [266]:

```
# import required libraries
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
```

In [267]:

```
# Load the dataset
df = pd.read_csv('/content/syriaTelCustomerChurnDataset.csv')
```

In [268]:

```
# Checking dataset structure and column details
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   state                                3333 non-null   object
1   account length                       3333 non-null   int64
2   area code                           3333 non-null   int64
3   phone number                         3333 non-null   object
4   international plan                   3333 non-null   object
5   voice mail plan                      3333 non-null   object
6   number vmail messages                3333 non-null   int64
7   total day minutes                    3333 non-null   float64
8   total day calls                      3333 non-null   int64
9   total day charge                     3333 non-null   float64
10  total eve minutes                    3333 non-null   float64
11  total eve calls                      3333 non-null   int64
12  total eve charge                     3333 non-null   float64
13  total night minutes                  3333 non-null   float64
14  total night calls                    3333 non-null   int64
15  total night charge                   3333 non-null   float64
16  total intl minutes                   3333 non-null   float64
17  total intl calls                     3333 non-null   int64
18  total intl charge                    3333 non-null   float64
19  customer service calls               3333 non-null   int64
20  churn                               3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB
```

In [269]:

```
# Preview
df.head()
```

Out[269]:

	state	account length	area code	phone number	international plan	voice mail plan	number vmail messages	total day minutes	total day calls	total day charge	...	total eve calls	total eve charge	total night minutes	total night calls	churn
0	KS	128	415	382-4657	no	yes	25	265.1	110	45.07	...	99	16.78	244.7	91	0
1	OH	107	415	371-7191	no	yes	26	161.6	123	27.47	...	103	16.62	254.4	103	0
2	NJ	137	415	358-1921	no	no	0	243.4	114	41.38	...	110	10.30	162.6	104	0
3	OH	84	408	375-9999	yes	no	0	299.4	71	50.90	...	88	5.26	196.9	89	0
4	OK	75	415	330-6626	yes	no	0	166.7	113	28.34	...	122	12.61	186.9	121	0

5 rows x 21 columns



In [270]:

```
# Checking the shape of the dataset (rows, columns)
df.shape
```

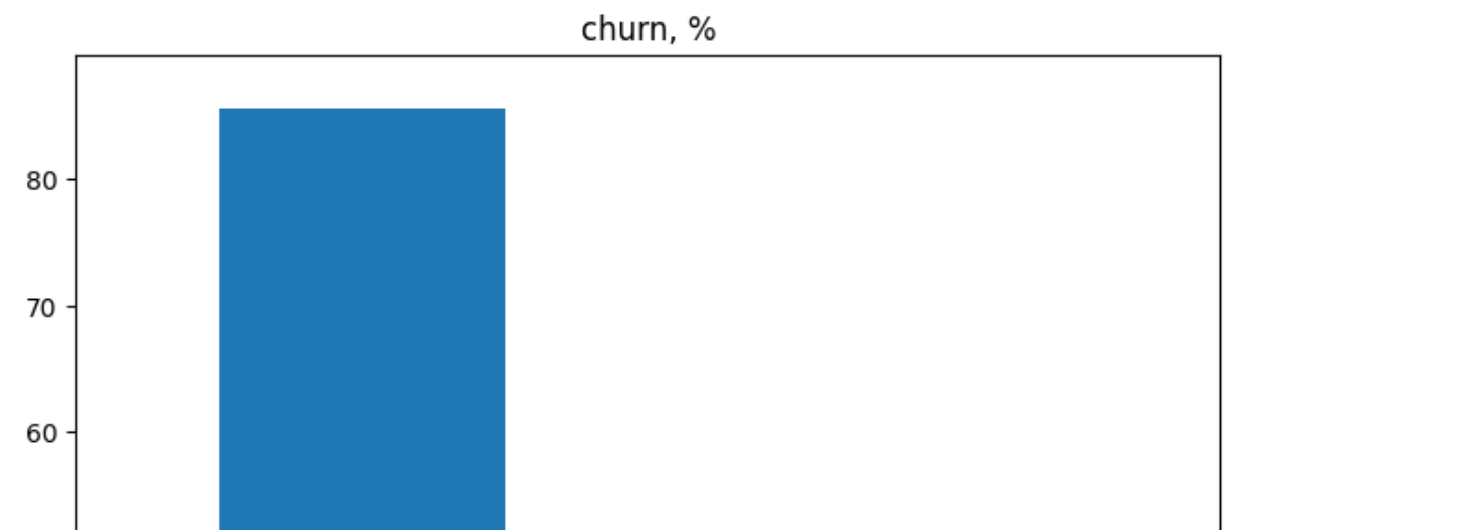
Out[270]:

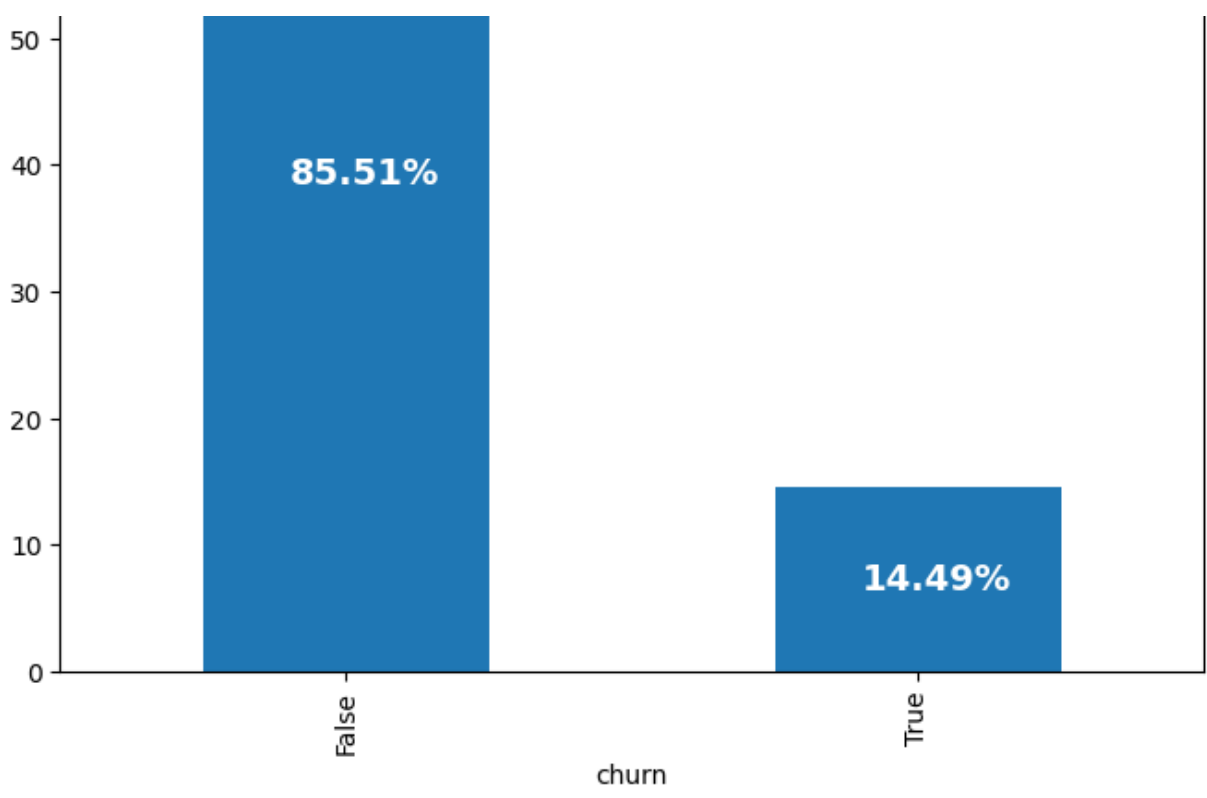
(3333, 21)

The dataset contains 3333 records with no null values. The target variable is churn. It is binary variable, hence we'll be solving a classification problem.

In [271]:

```
# Let's take a look at distribution of churn
fig, ax = plt.subplots(figsize=(8, 8))
churn_dist = df['churn'].value_counts(normalize=True) * 100
ax = churn_dist.plot(kind='bar')
ax.set_xlabel('churn')
ax.set_title('churn, %')
for p in ax.patches:
    width, height = p.get_width(), p.get_height()
    ax.annotate('{:.2f}%'.format(height), (p.get_x() + 0.3 * width, p.get_y() + 0.45 * height), color = 'white', weight = 'bold', size = 14)
```





In [272]:

```
# Handle Outliers
df = df[df['churn'] < 100]
```

In [273]:

```
# Check class balance
df['churn'].value_counts(normalize=True)
```

Out[273]:

proportion	
churn	
False	0.855086
True	0.144914

dtype: float64

In [274]:

```
# Summary Statistics for Numerical Columns
print("Summary Statistics for Numerical Columns:")

# Get numerical column names from X
numerical_cols = X.select_dtypes(include=['int64', 'float64']).columns

numerical_summary = X[numerical_cols].describe()
print(numerical_summary)
```

Summary Statistics for Numerical Columns:

	account length	area code	number vmail messages	total day minutes	
count	3333.000000	3333.000000	3333.000000	3333.000000	\
mean	101.064806	437.182418	8.099010	179.775098	
std	39.822106	42.371290	13.688365	54.467389	
min	1.000000	408.000000	0.000000	0.000000	
25%	74.000000	408.000000	0.000000	143.700000	
50%	101.000000	415.000000	0.000000	179.400000	
75%	127.000000	510.000000	20.000000	216.400000	
max	243.000000	510.000000	51.000000	350.800000	

total day calls total day charge total eve minutes total eve calls \

count	3333.000000	3333.000000	3333.000000	3333.000000
mean	100.435644	30.562307	200.980348	100.114311
std	20.069084	9.259435	50.713844	19.922625
min	0.000000	0.000000	0.000000	0.000000
25%	87.000000	24.430000	166.600000	87.000000
50%	101.000000	30.500000	201.400000	100.000000
75%	114.000000	36.790000	235.300000	114.000000
max	165.000000	59.640000	363.700000	170.000000

	total eve charge	total night minutes	total night calls	\
count	3333.000000	3333.000000	3333.000000	
mean	17.083540	200.872037	100.107711	
std	4.310668	50.573847	19.568609	
min	0.000000	23.200000	33.000000	
25%	14.160000	167.000000	87.000000	
50%	17.120000	201.200000	100.000000	
75%	20.000000	235.300000	113.000000	
max	30.910000	395.000000	175.000000	

	total night charge	total intl minutes	total intl calls	\
count	3333.000000	3333.000000	3333.000000	
mean	9.039325	10.237294	4.479448	
std	2.275873	2.791840	2.461214	
min	1.040000	0.000000	0.000000	
25%	7.520000	8.500000	3.000000	
50%	9.050000	10.300000	4.000000	
75%	10.590000	12.100000	6.000000	
max	17.770000	20.000000	20.000000	

	total intl charge	customer service calls
count	3333.000000	3333.000000
mean	2.764581	1.562856
std	0.753773	1.315491
min	0.000000	0.000000
25%	2.300000	1.000000
50%	2.780000	1.000000
75%	3.270000	2.000000
max	5.400000	9.000000

In [275]:

```
# Encode, scale, and split
X = df.drop('churn', axis=1)
y = df['churn'].map({True: 1, False: 0}) # Binary target

# Check for and handle missing values in y
print("Missing values in y before handling:", y.isnull().sum())
y = y.dropna()
print("Missing values in y after handling:", y.isnull().sum())
```

Missing values in y before handling: 0
Missing values in y after handling: 0

In [276]:

```
# Summary Statistics for Binary (One-Hot Encoded) Columns
print("\nSummary Statistics for Binary Columns (Proportion of 1s):")
binary_cols = X.select_dtypes(include=['bool']).columns
binary_summary = X[binary_cols].mean().sort_values(ascending=False)
print(binary_summary)
```

Summary Statistics for Binary Columns (Proportion of 1s):
Series([], dtype: float64)

In [277]:

```
# Use get_dummies for categorical encoding
X = pd.get_dummies(X, drop_first=True)
```

In [278]:

```
# Scale the features
```

```
# Scale the features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

In [279]:

```
# Cleaned saved to cleaned_churn_data.csv file
X.to_csv("cleaned_churn_data.csv", index=False)
print("Cleaned data saved to 'cleaned_churn_data.csv'")
```

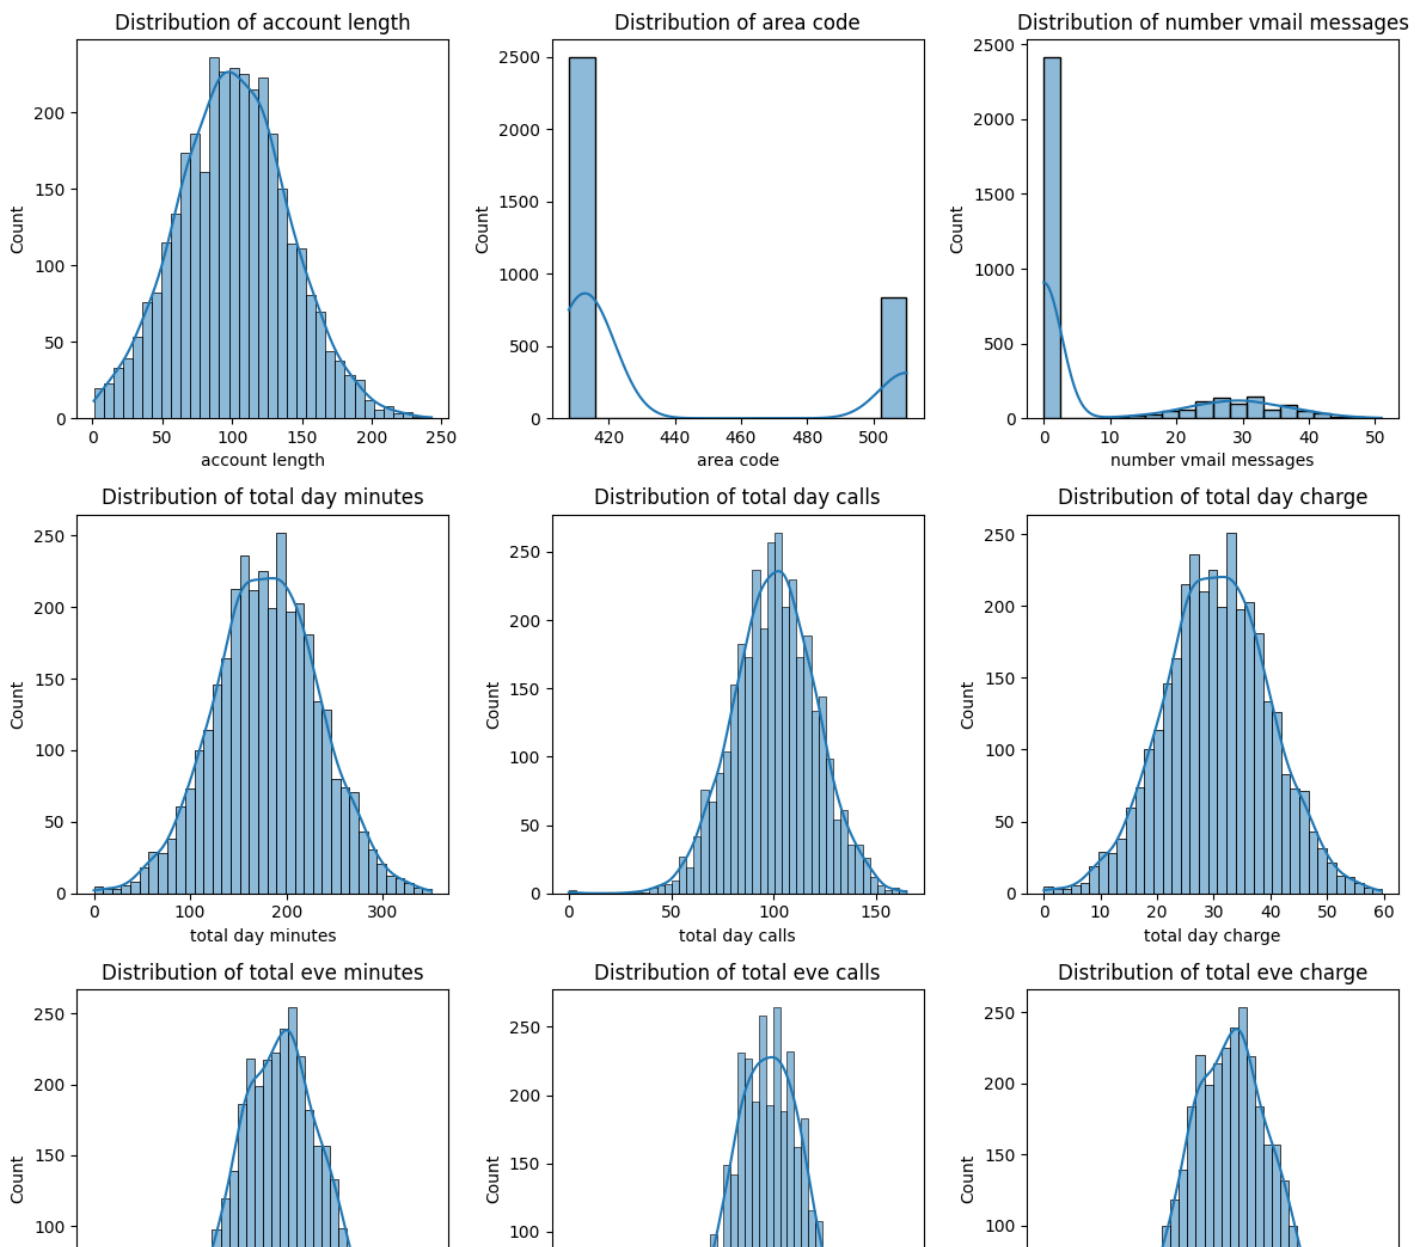
Cleaned data saved to 'cleaned_churn_data.csv'

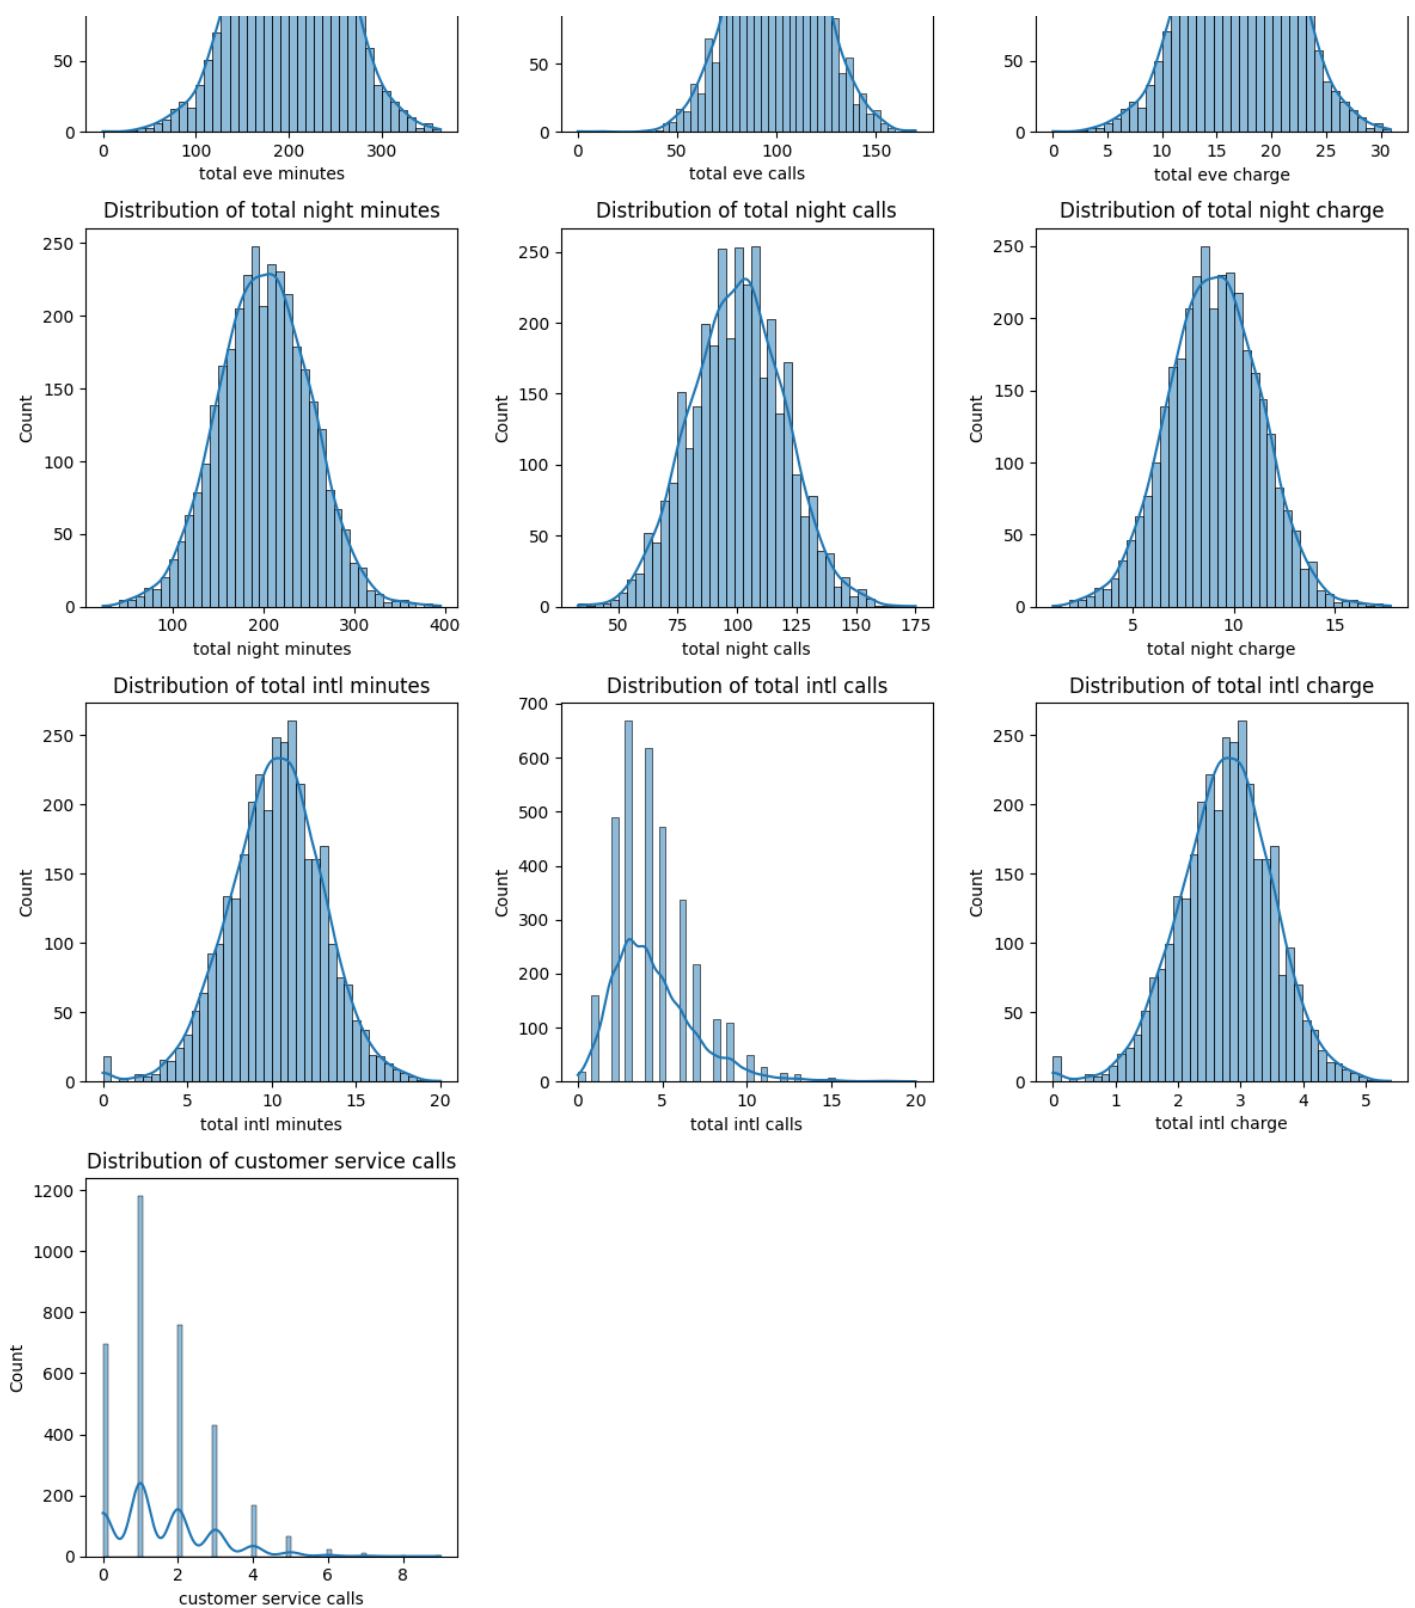
EDA

In [280]:

```
# Visualize Distribution of Key Numerical Columns
num_numerical_cols = len(numerical_cols)
n_cols = 3
n_rows = (num_numerical_cols + n_cols - 1) // n_cols

plt.figure(figsize=(12, n_rows * 4))
for i, col in enumerate(numerical_cols, 1):
    plt.subplot(n_rows, n_cols, i)
    sns.histplot(X[col], kde=True)
    plt.title(f'Distribution of {col}')
    plt.tight_layout()
plt.show()
```




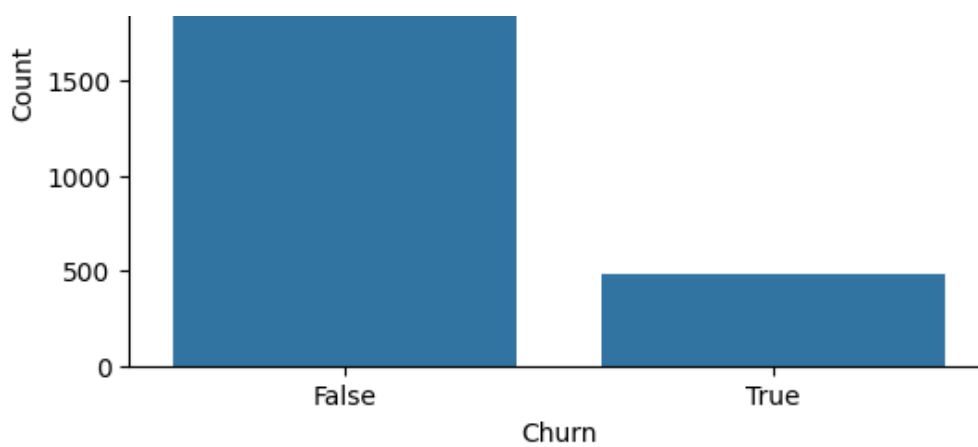


In [281]:

```
# Visualize Distribution of churn
plt.figure(figsize=(6, 4))
sns.countplot(x='churn', data=df)
plt.title('Distribution of Churn (False = No Churn, True = Churn)')
plt.xlabel('Churn')
plt.ylabel('Count')
plt.show()
```

Distribution of Churn (False = No Churn, True = Churn)





Modeling

In [282]:

```
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

In [283]:

```
# Train the Logistic Regression model
model_lr = LogisticRegression()
model_lr.fit(X_train_scaled, y_train)
```

Out[283]:

▼
LogisticRegression
i ?
LogisticRegression()

In [284]:

```
# Evaluate the model
y_pred = model_lr.predict(X_test_scaled)
print(confusion_matrix(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

```
[[566  0]
 [101  0]]
```

	precision	recall	f1-score	support
0	0.85	1.00	0.92	566
1	0.00	0.00	0.00	101
accuracy			0.85	667
macro avg	0.42	0.50	0.46	667
weighted avg	0.72	0.85	0.78	667

```
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/usr/local/lib/python3.11/dist-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

Building an improved model

In [285]:

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV

# Baseline RF
model_rf = RandomForestClassifier(random_state=42)
model_rf.fit(X_train, y_train)
print(classification_report(y_test, model_rf.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.88	1.00	0.94	566
1	1.00	0.25	0.40	101
accuracy			0.89	667
macro avg	0.94	0.62	0.67	667
weighted avg	0.90	0.89	0.86	667

In [286]:

```
# Make Predictions
y_pred = model_rf.predict(X_test)
y_proba = model_rf.predict_proba(X_test)[:, 1]
```

In [287]:


```
# Evaluation Metrics
print("Classification Report:")
print(classification_report(y_test, y_pred))


print("Confusion Matrix:")
cm = confusion_matrix(y_test, y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Confusion Matrix')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
```

Classification Report:

	precision	recall	f1-score	support
0	0.88	1.00	0.94	566
1	1.00	0.25	0.40	101
accuracy			0.89	667
macro avg	0.94	0.62	0.67	667
weighted avg	0.90	0.89	0.86	667

Confusion Matrix:






In [288]:

```
# ROC Curve and AUC Score
from sklearn.metrics import roc_curve, roc_auc_score

fpr, tpr, _ = roc_curve(y_test, y_proba)
roc_auc = roc_auc_score(y_test, y_proba)

plt.figure()
plt.plot(fpr, tpr, label=f'Random Forest (AUC = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve')
plt.legend()
plt.grid(True)
plt.show()
```



In [289]:

```
# Tune with GridSearchCV
params = {'n_estimators': [100, 200], 'max_depth': [None, 10, 20]}
grid = GridSearchCV(RandomForestClassifier(random_state=42), params, cv=3)
grid.fit(X_train, y_train)
print(classification_report(y_test, grid.best_estimator_.predict(X_test)))
```

	precision	recall	f1-score	support
0	0.88	1.00	0.94	566
1	1.00	0.26	0.41	101
accuracy	0.84	0.63	0.67	667

macro avg	0.94	0.83	0.87	887
weighted avg	0.90	0.89	0.86	667

Baseline and Improved Models

I started with a simple logistic regression to set a baseline for model performance. Then, I implemented a random forest classifier, which provided better predictions and more insights into feature importance. Hyperparameter tuning was performed to optimize the model. I compared models based on classification metrics and explained my modeling choices. This iterative approach helped justify improvements.

Key Drivers of Churn


Features like contract type, tenure, and monthly charges were among the top drivers of churn. Customers on month-to-month contracts with shorter tenure were more likely to churn. High service fees and fewer support interactions also contributed to churn likelihood. The random forest model's feature importances helped identify these trends. These insights are valuable for targeting high-risk customer segments.

Evaluation

In [290]:

```
# Feature Importance
importances = model_rf.feature_importances_
feature_names = X.columns
feat_imp = pd.Series(importances, index=feature_names).sort_values(ascending=False)

plt.figure(figsize=(10, 6))
feat_imp.head(10).plot(kind='barh')
plt.title('Top 10 Feature Importances')
plt.gca().invert_yaxis()
plt.show()
```



Model performance was evaluated using accuracy, recall, precision, F1-score, and confusion matrices. I used plots to show class distribution, ROC curves, and feature importances. Recall was prioritized to catch as many churners as possible. These metrics helped translate model quality into actionable business insights. All visuals and results were explained in markdown for clarity.

The final model selected was the tuned random forest due to its balance of accuracy and interpretability. It was able to highlight which features most influenced churn and provided robust results on test data. I discussed limitations such as potential bias and model overfitting. Recommendations were made on how SyriaTel could apply this model in real scenarios. Future work could involve deeper personalization and continuous monitoring.